

Algorithms and Complexity

Graphs: Representations and Exploration

Julian Mestre

School of Information Technologies
The University of Sydney



THE UNIVERSITY OF
SYDNEY


Undirected graphs

Let $G=(V,E)$ be an undirected graph:

- V = set of vertices (a.k.a. nodes)
- E = set of edges

Some notation

- $\deg(u)$ = # edges incident to u
- $\deg(G) = \max_u \deg(u)$
- $N(u)$ = neighborhood of u
- $n = |V|$
- $m = |E|$




Let $G=(V,E)$ be a directed graph:

- V = set of vertices (a.k.a. nodes)
- E = set of directed edges (a.k.a. arcs)

Some notation

- $\text{out-deg}(u)$ = # arcs out of u
- $\text{in-deg}(u)$ = # arcs into u
- $N^{\text{out}}(u)$ = out neighborhood of u
- $N^{\text{in}}(u)$ = in neighborhood of u




Graphs as a modeling tool

Can model many relations among elements in a set:

- Social network
- Internet topology
- Protein-protein interaction

Can help formulate problems:

- What's the distance between two nodes?
- What's a central node?
- How well connected the network is?
- What's a critical node?



Let $G=(V,E)$ be an undirected graph:


- Sequence v_1, v_2, \dots, v_k is a *path* if (v_i, v_{i+1}) is an edge in E for all $i=1, \dots, k-1$
- The *length* of the path is the # of edges in it
- A path is *simple* if no repeated vertices
- A *cycle* is a path v_1, v_2, \dots, v_k where $v_1 = v_k$
- A cycle v_1, v_2, \dots, v_k is *simple* if v_1, v_2, \dots, v_{k-1} is a simple path
- G is connected if every every vertex can reach every other vertex

We say that G is *tree* if:

- G is connected and doesn't have a cycle, or equivalently
- G is connected and $|E| = |V| - 1$

Exploring a graph from a starting vertex s , layer by layer:

- $L_0 = \{s\}$
- L_1 = vertices that are one hop away from s
- L_2 = vertices that are two hops away from s
- \vdots
- L_k = vertices that are k hops away from s



Breath first search



```
def BFS(G,s):  
  
    layers = []  
    current_layer = [s]  
    next_layer = []  
    "mark every vertex except s as not seen"  
    while "current_layer not empty" :  
        layers.append(current_layer)  
        for u in current_layer:  
            for v in "neighborhood of u":  
                if "haven't seen v yet":  
                    next_layer.append(v)  
                    "mark v as seen"  
        current_layer = next_layer  
        next_layer = []  
  
    return layers
```

Let G be a graph and s a vertex in G . Suppose we run $\text{BFS}(G,s)$ and that it returns layers L_0, L_1, \dots, L_k , then for u :

- If u belongs to some layer L_i , then there is a path from s to u
- If there is a path from s to u , then u belongs to some L_i
- In fact, u belongs to L_i if and only if the shortest s - u path has i edges

Edges across layers must connect adjacent layers.

Adjacency lists



Adjacency matrix

	u	v	x	y
u	0	1	0	1
v	1	0	1	1
x	0	1	0	0
y	1	1	0	0

Complexity depends on representation

Scan neighborhood of vertex **u**

- Adj. list : $\Theta(|N(u)|)$
- Adj. matrix : $\Theta(n)$

Check if **u** and **v** are adjacent:

- Adj. list : $\Theta(\min\{|N(u)|, |N(v)|\})$
- Adj. matrix : $\Theta(1)$

Space:

- Adj. list : $\Theta(|V|+|E|)$
- Adj. matrix : $\Theta(|V|^2)$

Time complexity of BFS

```
def BFS(G,s):
```

```
    layers = []
```

```
    current_layer = [s]
```

```
    next_layer = []
```

```
    “mark every vertex except s as not seen”
```

```
    while current_layer “not empty”:
```

```
        layers.append(current_layer)
```

```
        for u in current_layer:
```

```
            for v in “neighborhood of u”:
```

```
                if “haven’t seen v yet”:
```

```
                    next_layer.append(v)
```

```
                    “mark v as seen”
```

```
    current_layer = next_layer
```

```
    next_layer = []
```

```
    return layers
```

This takes
 $O(|V|)$ time

This loop takes
 $O(|N(u)|)$ time*

Adding up over all u , we get
 $O(\sum_u |N(u)|) = O(|E|)$

Graph:


- discrete object encoding a relation between vertices
- representations: adjacency lists, adjacency matrix
- time complexity of basic primitives depends on representation

Breath first search (BFS):

- a graph exploration strategy
- starting from a vertex s , visit vertices reachable from s , layer by layer
- L_i holds vertices at distance i from s
- Runs in $O(m+n)$ time if the graph is represented with adjacency lists


Depth first search

Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.



Depth first search


Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.



a ●



Depth first search

Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.





Depth first search

Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.





Depth first search

Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.





Depth first search

Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.





Depth first search

Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.





Depth first search

Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.





Depth first search

Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.



Depth first search

Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.



DFS tree

Depth first search

```
def DFS(G):
```

```
    “mark vertices as unvisited”
```

```
    “set vertices’ parent as None”
```

```
    time = 0
```

```
    for u in “vertices of G”:
```

```
        if “haven’t seen u yet”:
```

```
            DFS_visit(u)
```

```
    return parent
```

```
def DFS_visit(u):
```

```
    “mark u as visited”
```

```
    time = time + 1
```

```
    discovery[u] = time
```

```
    for v in “neighborhood of u”:
```


```
        if “haven’t seen v yet”:
```

```
            parent[v] = u
```

```
            DFS_visit(v)
```

```
    time = time + 1
```

```
    finish[u] = time
```



Time complexity of DFS

def DFS(G):

```
“mark vertices as unvisited”  
“set vertices’ parent as None”  
time = 0  
for u in “vertices of G”:  
    if “haven’t seen u yet”:  
        DFS_visit(u)  
return parent
```

Ignoring work done
inside function calls, it
runs in $O(n)$ time

def DFS_visit(u):

```
“mark u as visited”  
time = time + 1  
discovery[u] = time  
    for v in “neighborhood of u”:  
        if “haven’t seen v yet”:  
            parent[v] = u  
            DFS_visit(v)  
time = time + 1  
finish[u] = time
```

Ignoring work done
inside recursive calls,
it runs in $O(|N(u)|)$ time

$\Rightarrow O(m)$ time
overall here

The subset of edges $\{(u, \text{parent}[u]): u \in V\}$ forms a collection of trees (a.k.a. forest)

An undirected graph is connected if and only if we have a single tree in the DFS forest. In fact, each tree corresponds to a connected component of the graph.

Each discovery and finishing time is a unique number in $[1, 2n]$

Def.: In a connected graph $G=(V,E)$, we say that $(u,v) \in E$ is a *cut edge* if $(V, E-(u,v))$ is not connected.

Trivial algorithm runs in $O(m^2)$: For each edge $(u,v) \in E$

- Remove (u,v) from G
- Run DFS to check if the new graph is still connected

A better algorithms runs in $O(nm)$: For each edge in DFS tree



- Remove (u,v) from G
- Run DFS to check if the new graph is still connected

Back edges in DFS forest

Def.: If $\text{parent}[u] \neq \text{None}$ then we call $(u, \text{parent}(u))$ a *tree edge*

Def.: We say that a non-tree edge (u, v) is a *back edge* if u is a descendant of v in the DFS forest, or vice-versa

Obs.: In the DFS forest every non-tree edge is a back edge



Def.: In a connected graph $G=(V,E)$, we say that $u \in V$ is a *cut vertex* if $G-u$ is not connected.

Obs.: If u is the root of the DFS tree, then u is a cut vertex if and only if it has two or more children

Obs.: If u is a leaf of the DFS tree, then u is not a cut vertex

Obs.: If u is not the root of the DFS, u is a cut vertex if it has a child v and no vertex in v 's subtree can “jump over” u

Let u be an internal vertex in the DFS tree.

Def.: $up[u] = \min discovery[v]$ where $v \in N(u)$

v is allowed
to be u

Def.: $down\&up[u] = \min up[v]$ where v is a descendent of u

An internal vertex u is a cut vertex if and only if it has a child v :
 $down\&up[v] = discovery[u]$

Only need to show how to compute $down\&up$ in $O(m)$ time

Thm.

Given a connect graph, there is an $O(m)$ time algorithm for computing its cut vertices

Another graph exploration strategy: Follow edges to new nodes until “stuck”, then backtrack

Vertices are assigned a discovery and a finishing time

In undirected graph, each edge can be a tree edge or a back edge

DFS is useful for solving other graph problems, like cut edges and cut vertices

Runs in $O(m+n)$ time if graph is represented with adjacency lists


Simple Python graph representation

Let $G=(V, E)$ be a graph on n vertices, then:

- Vertices can be any non-mutable object: e.g., a string, or a number
- G is a dictionary (dict)
- $G[u]$ = adjacency list of vertex u

Example:

```
-G = { "a": ["b", "c"],  
      "b": ["a", "d", "c"],  
      "c": ["a", "b"],  
      "d": ["b"] }
```



- To iterate over vertices in G use \longrightarrow `for u in G:`
- To iterate over neighbors of “a” use \longrightarrow `for u in G[“a”]:`

Quiz 2

- 15 minutes long
- during your tutorial session

Problem set 2:

- will be posted later today (Monday 6 August)
- make sure you work on it before the tutorial

Assignment 1:

- Is due tomorrow (Tuesday 7 August)

Assignment 2:

- Out tomorrow, due next Tuesday